



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Counterexamples for a Rust Verifier

Bachelor Thesis

Cedric Hegglin

January 19, 2038

Advisors: Prof. Dr. Peter Müller, Dr. Christoph Matheja, Aurel Bily

Department of Computer Science, ETH Zürich

Abstract

Contents

Contents	iii
1 Introduction	1
2 Overview	3
2.1 Counterexamples	3
2.2 Verification Pipeline	5
3 Approach	9
4 Implementation	13
4.1 Improving Counterexamples in Silicon	13
4.1.1 Native Models	13
4.1.2 ExtractedModel	15
4.1.3 Evaluating Terms	16
4.1.4 Heap Extraction	18
4.1.5 Final Mapping	18
4.2 Prusti Counterexamples	21
4.2.1 Identifying the failing function	21
4.2.2 Variables	21
4.2.3 Backtranslation	22
4.2.4 Choosing Labels	27
4.2.5 Presentation	30
5 Evaluation	33
5.1 Qualitative Evaluation of Counterexamples	33
5.2 Timing Analysis	41
6 Conclusion	43
A Appendix	45

CONTENTS

Bibliography

47

Chapter 1

Introduction

Software is all around us and as our dependency on it is constantly growing, ensuring its correctness is becoming ever more important. Unfortunately, especially in the context of systems programming, with conventional languages like C or C++, writing correct programs is notoriously difficult. The Rust programming language [2] is an alternative to those languages, prohibiting programmers from introducing some of the most common errors like data-races by enforcing a strict type system and ownership model.

Prusti [4] is a formal verifier for the Rust programming language, extending Rust's guarantee of memory-safety to functional correctness. It leverages those guarantees to allow users to add specifications to Rust programs and prove them without requiring the level of expertise that is usually needed to verify properties in verifiers for other system programming languages. This makes Prusti a very accessible tool for a broad range of developers. However, verifying programs is still not an easy task, and especially when verification fails, the information Prusti (in its state before this project) provided to the user was very limited. Counterexamples, i.e. the values of the variables that make an error reproducible, could be of great help in this situation and further improve the usability of Prusti. To give an example, let us look at the following Rust function:

```
#[ensures(result == n*(n+1)/2)]
fn sum(n: i32) -> i32 {
    if n <= 0 {
        0
    } else {
        sum(n-1) + n
    }
}
```

For positive n this function computes $\sum_{i=1}^n i$ which is indeed equal to $\frac{n*(n+1)}{2}$, but negative inputs will result in a violation of the post-condition. However, if a Prusti user is not aware of this fact, in the initial state of Prusti all the information one could get is that verification fails and that the post-condition might be violated. A counterexample for this function would for example provide the values `n <- -1`, `result <- 0` and explains the reason for the verification error. Counterexamples are a highly useful feature of formal verifiers which is why we set the goal to extend Prusti with this feature.

Prusti is one of several verifiers built as a front-end of the Viper infrastructure [8], using the verifier Silicon. Prusti translates Rust programs to the Viper intermediate verification language and uses Silicon's verification result to reason about Rust programs. Since Silicon already had some support for counterexamples, the main approach to make this project feasible was to use these existing Silicon counterexamples and translate them back to meaningful counterexamples for Rust programs.

However, Silicon's counterexamples initially contained various Viper internal representations and its information was spread over various data structures. Therefore, part of the effort of this project went into making Silicon's representation of counterexamples more accessible and compact. This leads us to our main contributions:

- we improved the counterexample support of Silicon for both direct users and front-ends by adding an additional simplified representation.
- We implemented the translation of said Counterexamples back to their corresponding Rust values for various data types.
- We discussed presentation of counterexamples on various examples and implemented it for Prusti.
- We discussed the current limitations of our implementation and its source and explored some approaches on how to resolve them in the future.

Chapter 2

Overview

In this chapter we will cover some details on counterexamples and also have a look at the parts of Prusti and Silicon that are relevant for this thesis in their state before this project.

2.1 Counterexamples

When trying to verify a function, it can be challenging to determine the source of a verification error. If we run Prusti in its state before our project on the function given in listing 2.1, it only reports that its post-condition might not hold as in listing 2.2. The function computes the maximum of two values and we try to prove that its result has to be larger than one of the two values. Prusti fails to verify this function because it violates the post-condition if the two values are equal. While it might be easy to figure out why the verification fails in this case, if one fails to do so, Prusti provides no further support.

Listing 2.1: Example Rust function with Prusti annotations

```
#[ensures(result > x || result > y)]
fn max(x: i32, y: i32) -> i32 {
    if x > y {
        x
    } else {
        y
    }
}
```

Ideally, Prusti should not only tell us that verification failed but justify the failure with a counterexample. As already mentioned in chapter 1, counterexamples define execution traces of the program that lead to an error

Listing 2.2: Output generated by Prusti in its initial state

```
error: [Prusti: verification error] postcondition might not hold.
--> counterexample-thesis-resources/examples-report/prusti/max.rs:4:11
|
4 | #[ensures(result > x || result > y)]
|           ~~~~~
|
...
Verification failed
error: aborting due to previous error; 1 warning emitted
```

state or violation of the specification. In the context of SMT-based verifiers, counterexamples usually provide a model containing the values of variables at various program points. We speak of a *valid counterexample* if it does not violate any pre-condition, matches the values of an actual execution of the program and leads to some error state. If verification fails proving the correctness of a certain program, this can give us valuable information why it is failing. For the previous example $\{x:2, y:2\}$ is a valid counterexample because there are no pre-conditions and executing the function with these arguments will always lead to a violation of the post-condition, whereas $\{x:1, y:2\}$ is not.

To give an idea of different ways to present counterexamples we are looking at two existing verifiers that already implement counterexamples. One of those verifiers is Nagini [6], another Viper frontend but for the Python programming language. If we run Nagini on the analogous Python function of our previously considered example, we get the result presented in listing 2.3. It provides the values of the arguments when entering the function and the values of all local variables when the function returns. The value False for x might be confusing, but it stems from the fact that booleans are a subtype of integers in Python and represent zero as well. Therefore this is a valid counterexample.

Listing 2.3: Counterexample given by Nagini

```
Old store:
  x -> False,
  y -> 0
Old heap: Empty.
Current store:
  x -> False,
  y -> 0,
  Result() -> 0
Current heap: Empty.
```

```

method Max(x:int, y:int) returns (n:int)
  ensures (n > x || n > y)
  {
    x = 0; y = 0;
    if (x>y){
      n := x;
    } else {
      n := y;
    }
  }

```

Figure 2.1: Method annotated with Dafny's counterexample

Another widely used Verifier is Dafny [7]. When using Dafny with its Visual Studio Code (VSCode) Extension, it can display the values of the counterexample directly within the code (in blue), showing the values of the local variables at the various points during execution and therefore also displaying which paths are taken. Its result for the same function as before is shown in fig. 2.1. Additionally, even though it is not displayed in the given example, Dafny also displays the values of variables at each point where their values change and even displays the return values of function calls within the method. This can be crucial in some cases, for example if there is a call to a non deterministic function like a random number generator and the verification error depends on its return value. It can potentially make it very hard to reproduce or understand the error without this information.

On a more general note on counterexamples in the context of symbolic execution, we should know they can be *spurious*, or the underlying SMT-solver can fail to generate any counterexample. Spurious counterexamples can occur due to the abstractions performed during symbolic execution since it is usually sound but not complete. But even with this limitation, getting a spurious counterexample can be more useful than not getting any information because it tells us that this proof fails because of the limitations of the verifier and not necessarily because of an implementation error.

2.2 Verification Pipeline

To be able to understand the approach and implementation of this project, a basic understanding of Prusti and Silicon in their state before our contribution is necessary. In this section we will cover the relevant parts of Prusti's verification process and Silicon's available but unused counterexamples by going through various parts of the verification process for our previous example from 2.1.

It is important to know that Prusti is an extension of the Rust compiler and has access to most of its internal utilities. Given a Rust program, Prusti invokes the Rust compiler on this program and mostly works with the mid-

level intermediate representation (MIR) of the program. In this representation the result of a function will always be stored in the variable `_0` and the arguments are the subsequent variables (`_1` and `_2` in our example). Prusti then proceeds to translate the generated MIR encoding to a Viper program. This translation is quite complex but in the context of this project we are mostly focusing on the encoding of various types. Chapter 4 will delve more into this translation, but the resulting arguments of the generated program for our considered example will have the following structure:

```
_1 <- Ref ($Ref!val!0) {  
    val_int: int  
}  
_2 <- Ref ($Ref!val!1) {  
    val_int: int  
}
```

After the translation Prusti runs Silicon to verify the generated Viper program and uses its result. If the verification fails, Silicon's errors are mapped to their origin in the Rust program. In this case, Silicon is able to produce a counterexample but Prusti did not take advantage of that capability. Since this project heavily relies on the counterexamples provided by Silicon, we need to understand them in their state before our contribution. Their implementation is the result of a previous thesis called "SMT Models for Verification Debugging" [9] and in the following we will have a first look at what information they provide.

In the initial state of Silicon, it generates a so called native counterexample and a state for each method that fails verification. The native counterexample is a model generated by the SMT solver. It is a map of identifiers to the model's entries which can be values of variables at certain points in the program or results of called functions among other things. The state contains various data structures of which the following are relevant:

- the store, a map from the names of local variables to their value or another identifier pointing to some location in the other data structures;
- the main heap, containing so called chunks of predicates and fields of references representing their values when the method fails verification; and
- the old heaps, a map from labels to heaps, where each heap has the same structure as the main heap, but it represents all the values of the method's variables when crossing the corresponding label. It always contains a heap at the label "old", representing the values when entering the method, and additional heaps if the execution runs through labels that are specified within the methods body.

To get an idea of how we can use these counterexamples but also to emphasize how cumbersome it was to use them in their initial state, we will look at some concrete examples. First we consider a heap-independent Viper method (listing 2.4) that tries to prove the same thing as our previous example. A valid counterexample would give us three equal values for both of the arguments and the result. If we look at the relevant parts of the generated counterexample (listing 2.5) it is fairly easy to extract this information by mapping the provided identifiers from the store to the given model entries, giving us a value of 906 for all the variables.

```
method max(x: Int, y: Int) returns (n: Int)
  ensures(n > x || n > y)
  {
    if (x > y) {
      n := x
    } else {
      n := y
    }
  }
```

Listing 2.4: Heap independent Viper example

```
Store {
  x -> x@3@04,
  y -> y@4@04,
  n -> y@4@04,
}

Model {
  x@3@04 -> 906,
  y@4@04 -> 906,
}
```

Listing 2.5: relevant parts of counterexample for listing 2.1

Unfortunately, extracting this information quickly becomes a lot more complicated once we start dealing with values involving the heap. To demonstrate the problem we will have a look at a simplified version of the method that is generated by Prusti for our initial example (listing 2.6). We do not use Prusti's generated Viper code because it is hard to read and the produced files are a lot larger. The fields of the arguments are now located on the heap and hence, the resulting counterexample suddenly involves a lot of Viper internal structures. Extracting meaningful information, even with the required knowledge, is tedious. If we wanted to get the value of `_1`, we would first have to get the identifier `_1@3@04` from the store. Then on the heap we can find the location of the field for `_1.val_int` given by this `SortWrapper` construct, which we do not have to understand completely yet,

but it points us to the matching field of the model where we can get a value of 323 for this field. This is still a fairly simple example, but the point to be made here is that the current interface for counterexamples in Silicon is hard to use for both direct users and front ends of Silicon.

Listing 2.6: Viper program similar to the one generated by Prusti

```
field val_int: Int
method max(_1:Ref, _2:Ref) returns (_0:Int)
  requires(acc(_1.val_int) && acc(_2.val_int))
  ensures(_0 > old(_1.val_int) || _0 > old(_2.val_int))
  {
    if (_1.val_int > _2.val_int) {
      _0 := _1.val_int
    } else {
      _0 := _2.val_int
    }
  }
}
```

Listing 2.7: relevant parts of counterexample for listing 2.6

```
Store {
  _1 -> _1@3@04,
  _2 -> _2@4@04,
  _0 -> _0@9@04,
},
Heap{
  _2@4@04.val_int -> SortWrapper.$SnapToInt(Second:($t@6@04)),
  _1@3@04.val_int -> SortWrapper.$SnapToInt(First:($t@6@04)),
},
Model {
  $t@6@04 -> ($Snap.combine $Snap.unit $Snap.unit),
  $SortWrappers.$SnapToInt -> {
    323
  },
  _0@9@04 -> 323,
}
```

Chapter 3

Approach

As outlined in chapter 1, the idea of this project is to use the counterexamples generated by Silicon to generate meaningful counterexamples for Rust programs. On a high level, achieving this goal can be split into three parts: extracting the needed information from Silicon's counterexamples, translating this information back to their corresponding Rust components and finally, presenting the counterexamples to the user.

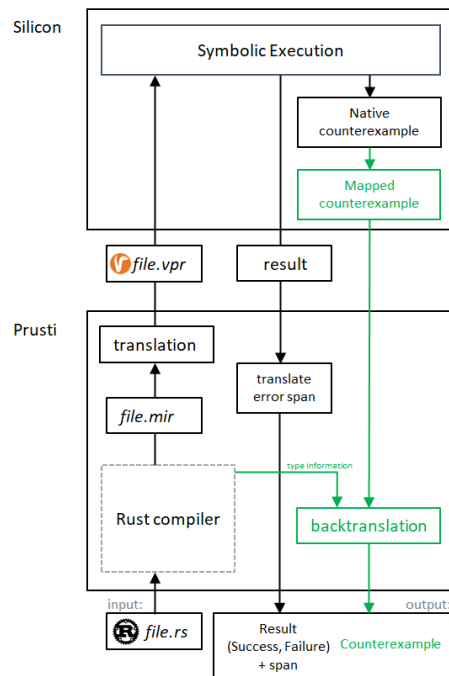


Figure 3.1: Prusti's structure with new components (green).

3. APPROACH

As we have seen previously, the counterexamples that are initially available in Silicon are difficult to work with and if we are interested in the values of fields of references, there is a lot of processing to be done. Since almost every Rust type is translated to a reference in Viper, we added a simpler representation of counterexamples to Silicon that contains an entry for each variable. Such an entry recursively contains other entries for its fields if it is a reference. To obtain this new representation it evaluates all Silicon internal values found on the heap and within the model that we have seen previously and maps the given values of the fields to the corresponding references. This is also done with the heaps at all the old labels that are available, giving us multiple values for variables that are mutated during execution. While all of this could be implemented within Prusti, we decided to do it on the Silicon level and design it in a way that makes it front-end independent. This has the advantage that it might be used in future implementations of counterexamples for other front-ends but can also be useful to Viper users in general.

To give an idea of what our new representation achieves, we take a look at another version of our previous example shown in listing 3.1, computing the maximum of two values with the slight difference that the permissions to the fields are inhaled in the methods body instead giving them as a precondition. Listing listing 3.2 shows us the new representation of the counterexample that is generated for the example. Looking at the values given when the method returns, we can clearly see the advantage over the initial native counterexamples. However it is worth noting that at the label "old", where the permission to the fields is not inhaled yet, we do not get any values for the fields of the references.

Listing 3.1: adjusted Viper example

```
field val_int: Int
method max(_1:Ref, _2:Ref) returns (_0:Int)
  ensures(acc(_1.val_int) && acc(_2.val_int))
  ensures(_0 > _1.val_int || _0 > _2.val_int)
  {
    inhale(acc(_1.val_int) && acc(_2.val_int))
    if (_1.val_int > _2.val_int) {
      _0 := _1.val_int
    } else {
      _0 := _2.val_int
    }
  }
```

Listing 3.2: mapped counterexample for previous example

```
model at label: old
_1 <- Ref ($Ref!val!0) {
}
_2 <- Ref ($Ref!val!2) {
}
_0 <- 906

on return:
_1 <- Ref ($Ref!val!0) {
    val_int(perm: 1/1) <- 906
}
_2 <- Ref ($Ref!val!2) {
    val_int(perm: 1/1) <- 906
}
_0 <- 906
```

On the Prusti side, we now use Silicon’s new representation combined with the compiler’s information to obtain counterexamples for Rust. We first need to find a mapping of all Rust variables of the failing function to their corresponding variable names in Viper. Then we try to get a counterexample entry for each Viper variable and invert the translation of the various types that Prusti performed, according to the type of the Rust variable which we get from the compiler. We try to translate every variable according to its type if the necessary information is available, which needs to be done recursively for types like tuples, structs or pointers. For example, if we assumed that `_1` was the entry for `x` back in our example listing 2.1, the compiler would tell us that `x` is an integer and we would get its value from `_1.val_int`. This sort of backtranslation is supported for integers, booleans, characters, references, tuples, structs and enums.

The last part of our approach is presenting the counterexamples to the user. Before even starting any implementation we already used a large set of Rust examples that fail verification to discuss how to meaningfully present counterexamples for them. One way of presenting them is to give the initial value of all arguments, and all the values of the local variables including the result and the arguments once the function fails. In almost all considered cases this makes determining the reason for the verification error very straightforward. That’s why, for any failing function, Prusti now computes the available values at those two points. They are presented to the user as a compiler note for each variable. This serves for compatibility with Prusti-Assistant, the Visual Studio Code extension, and in the future this might be used to see the failing values for a variable by simply hovering over it in the editor.

3. APPROACH

```
note: counterexample for "x"
initial: x <- -145
final: x <- -145
--> counterexample-thesis-resources/examples-report/prusti/max.rs:5:8
5 | fn max(x: i32, y: i32) -> i32 {
  |               ^
```

Figure 3.2: Prusti's output for variable x for listing 2.1

Chapter 4

Implementation

In this chapter we will delve into the details of the implementation of our project. It will be split into two parts, first explaining the changes we applied to Silicon to obtain simpler, more compact representation of counterexamples and secondly explaining how this new representation can be used in Prusti to generate counterexamples for Rust programs.

4.1 Improving Counterexamples in Silicon

In this section we go into more detail on the implementation of the new mapped Silicon counterexamples. Some parts of the implementation are based on Nagini’s [6] implementation, which is the first Viper front-end that implemented counterexamples. With the goal of making the available information more usable, we implemented a new class called `MappedCounterexample` and a class `Converter` that is responsible for translating the existing information to the new representation. The new counterexamples are generated if the user executes Silicon with the flag `“counterexample --mapped”`.

4.1.1 Native Models

First we need to have a deeper look at the initially available native model that we already briefly discussed in Chapter 2. While the heap and store are also crucially important to process the model correctly, all of the information of a counterexample that tells us when a method fails is found in the model. The native model is a map of strings to instances of a class called `ModelEntry`. The structure of this class and its sub-classes is shown in Figure 4.1.

Instances of the class `ConstantEntry` represent, as the name suggests, constant values given as strings. These can be the strings of integer or boolean values, strings of identifiers (e.g. for references) or the string of internal

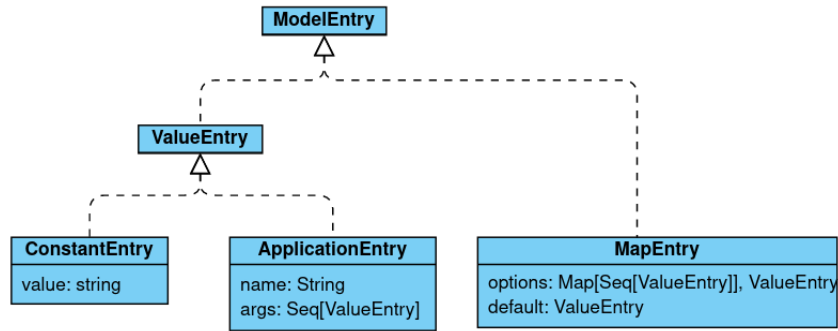


Figure 4.1: UML diagram of the ModelEntry class

values such as `Snap.unit`. The class `ApplicationEntry`, which is also part of the class `ValueEntry`, contains the name of a function and a sequence of arguments. The following listing contains examples of application-entries we typically encounter in generated models:

```

$t@2@04 -> ($Snap.combine $Snap.unit $Snap.unit)
x@1@04 -> (- 563)

```

In the printed output of counterexamples, shown in the above listing, application-entries are always displayed as S-expressions meaning that they are put into parentheses containing their function name followed by the arguments, which are either constant-entries or other application-entries, allowing for a nested structure.

Instances of the class `MapEntry` represent the result of function calls and define a mapping from a sequence of value entries to a single value-entry. They store results of function calls during the execution and help to recover those results for given values of the arguments. Those entries are not only generated for functions that are explicitly defined in the program we are verifying, but also for Silicon internal functions like sort-wrappers. To get a result for a given sequence of arguments we try to find the arguments in the map options (Figure 4.1) and use the value-entry it maps to if we find it. Otherwise we take the default value, which is shown as the `else` case in the printed output. In the printed version of the model map-entries are always surrounded by curly brackets as we can see in the following listing:

```

SortWrappers.SnapToInt -> {
  (Snap.combine Snap.unit Snap.unit) -> 21
  else -> 42
}

```

4.1.2 ExtractedModel

Before looking into the translation and processing of the heap and state using the native model, we need to know what our desired result will look like. The new model will consist of several models containing the values of local variables at each label of a method. Each model is a map from identifiers of variables to a new class of entries. The new class is called `ExtractedModelEntry` and contains sub-classes for the various values we are interested in for a meaningful counterexample. Listing 4.1 shows the various types of entries our new model will contain. The first three classes represent values of Viper’s basic types. Then there are three classes to represent references. In Silicon every instance of a reference has a unique identifier of the form `Ref!Val!0` which is stored in the `name` field of these entries. The class `RefEntry` also contains all of its fields as a mapping from the field name to

Listing 4.1: structure of `ExtractedModelEntry`

```
sealed trait ExtractedModelEntry

//Literals
case class LitIntEntry(value: BigInt)
case class LitBoolEntry(value: Boolean)
case class LitPermEntry(value: Double)

//References
case class RefEntry(
  name: String,
  fields: Map[String, (ExtractedModelEntry, Permission)]
)
case class NullRefEntry(name: String)
case class RecursiveRefEntry(name: String)

//Sequences
case class SeqEntry(
  name: String,
  values: Vector[ExtractedModelEntry]
)

//Debugging and intermediate results
case class VarEntry(value: String, sort: Sort)
case class OtherEntry(value: String, problem: String)
case class UnprocessedModelEntry(entry: ValueEntry)
```

another entry and the permission to that field. The class `RecursiveRefEntry` was added to avoid infinite recursion for cyclic references. We also added support for sequences while other composite types like sets and multi-sets remain as future work. Finally there are a few types of entries that either store some intermediate result or help reporting the reason why a value of an entry could not be extracted. For example instances of the class `VarEntry` store identifiers of sequences and references that will be replaced with their corresponding entries at some point.

4.1.3 Evaluating Terms

Both on the heap and in the store, we will encounter instances of Silicon's `Term` class. To get a meaningful counterexample we need to be able to evaluate those terms for a given native model. We implemented a function `evaluateTerm` which takes a native model and a term as arguments and returns an `ExtractedModelEntry`. Its implementation is based on Nagini [6], with the difference that Nagini's version operates on an older version of the native model and that its method only returns strings. Our function supports a subset of Silicon's terms and in the following we are going to explain how the most frequently occurring terms are evaluated.

Variables

Variable terms are of the form `Var(name, sort)` where the sort defines their type. To evaluate them, we always try to find an entry in the native model mapped to their name. If we do not find one we return an `OtherEntry` reporting that it can not be found. Otherwise, what we do to evaluate the found entry depends on the sort of the variable; there are three cases:

- If the sort is `Int`, `Bool` or `Perm`, we try to get an actual value from the native entry and return a literal entry. The found entry either has to be a constant-entry or an application-entry of the function `"-"` containing another constant, meaning that the contained value has to be negated. The constant-entry can be parsed to a value of the required type. The function, if there is one, has to be applied to get the final value. We return one of the three literal entries depending on the sort. Note that the sort `Perm` is not only used for fractional permissions but also for the type `Rational` in Viper. The following listing contains three examples of native entries for each of the discussed sorts:

```
x@0@04 -> 1.0,  
y@1@04 -> False,  
z@2@04 -> (- 7)
```

- If the sort is `Ref` or `Sequence`, then the found entry needs to be a constant containing the string of an identifier. The evaluation returns

a variable-entry containing the found identifier and its sort. We are interested in the values of the sequence or the fields of the reference, but they have to be resolved at a later stage, and the variable-entry serves as a placeholder. The following listing contains native entries generated for a sequence and a reference:

```
x@1@04 -> Seq<Int>!val!1,
y@4@04 -> $Ref!val!2
```

- If the sort is Snap, we return an `UnprocessedModelEntry` directly containing the found native entry. Why this is needed will become clear when we evaluate the next term.

First/Second

The terms `First` and `Second` are part of Silicon's modelling of the heap. They can only be applied to application-entries of the form `(Snap.combine x y)` where `First` returns `x` and `Second` returns `y`. They contain another term that, when evaluated recursively, should return an `UnprocessedModelEntry` containing the entry to which we can apply them. As an example we want to evaluate `Second(t@4@04)` with the given model entry:

```
t@4@04 -> (Snap.combine Snap.unit (Snap.combine Snap.unit Snap.unit))
```

We first recursively evaluate the variable `t@4@04`, returning an unprocessed model-entry. We then apply the function left or right to the contained entry, giving us the application-entry `(Snap.combine Snap.unit Snap.unit)` which we return in another unprocessed model-entry.

Sort-Wrappers

The last terms we are covering are the sort-wrappers. They contain two sorts and another term. The two sorts are needed to construct the name of the map-entry it refers to in the model, e.g. `SnapToRef`. To evaluate a term, for example `SortWrapper.SnapToRef(First(Second(t@4@04)))`, we first evaluate the contained term recursively. It should return another unprocessed model-entry containing an application-entry, for example `(Snap.combine Snap.unit Snap.unit)`. In the native model we should find a `MapEntry` of the following form:

```
SortWrapper.SnapToRef -> {
  (Snap.combine Snap.unit Snap.unit) -> Ref!val!4
  else -> Ref!val!3
}
```

The returned value of our recursive evaluation maps to `Ref!val!4`, and therefore we return this identifier in a `VarEntry`. Depending on the sort we expect, the result could also be one of the literal entries.

4.1.4 Heap Extraction

With the new function `evaluateTerm`, we can now start processing the heaps at the various labels. A heap in Silicon is a list of chunks. There are multiple types of chunks in Silicon but we only handle so-called basic-chunks which are either predicate or field-chunks. This is another design decision we reproduced from Nagini [6]. A field-chunk consists of an argument term, a field name and a snap term and tells us that the value of `arg.fieldname` is given by the snap term. For example, if we run Silicon on the method in Listing 4.3, it generates the main heap and the heap at the label "old". Since no values change during the execution of this method, we only need to look at one of the heaps, which is shown in the following:

```
SortWrapper.SnapToRef(First(t@2@04).field1 ->
    SortWrapper.SnapToInt(Second(Second(t@2@04))),
x@1@04.next -> SortWrapper.SnapToRef(First($t@2@04)),
SortWrapper.SnapToRef(First(t@2@04).field2 ->
    SortWrapper.SnapToInt(First(Second(t@2@04))),
```

The goal of the extraction is to generate an expression of the form `identifier.fieldname = entry` for each field-chunk. We can get that information by using our previously discussed function, `evaluateTerm`, on both the argument and the snap-term. The argument has to evaluate to a `VarEntry` and we only keep its name. For the above example the result for the extraction, given the native model in listing 4.4, is the following:

Listing 4.2: extracted heap

```
Ref!val!2.field1 <- 1,
Ref!val!0.next <- Ref!val!2,
Ref!val!2.field2 <- 0
```

All the extracted information is stored in a new `ExtractedHeap` class and will be used for the final mapping as explained in the next section. While predicate chunks might also be useful for some applications, we currently are not using them for our counterexamples in Prusti. They are also processed and part of the extracted heap but for the final mapping only field chunks are relevant, which is why we are not covering them in more detail.

4.1.5 Final Mapping

After processing all the heaps we can now try to get the actual values of a counterexample. We evaluate the store at each of the given heaps to get

Listing 4.3: Example Viper program with nested references

```
field field1: Int
field field2: Int
field next: Ref

method simple(x: Ref)
  requires acc(x.next)
  requires acc(x.next.field2) && acc(x.next.field1)
  {
    assert(x.next.field2 == x.next.field1)
  }
```

Listing 4.4: native counterexample for listing 4.3

```
x@1@04 -> Ref!val!0,
t@1@04 -> (Snap.combine
  Snap.unit
  (Snap.combine
    Snap.unit
    (Snap.combine Snap.unit Snap.unit)
  )
),
SortWrappers.SnapToInt -> {
  (Snap.combine Snap.unit Snap.unit) -> 1
  else -> 0
},
SortWrappers.SnapToRef -> {
  Ref!val!2
},
```

the values of the method at those labels. Since no values change during the execution of the previous example, we will again only look at one heap. Now our goal is to get a counterexample value for each variable that is located on the store. The store is a mapping of variables to terms and for our previous example in Listing 4.3 it only contains the value $x \rightarrow x@1@04$.

If we want to get an entry for the variable x , we start by once again using our `evaluateTerm` function on $x@1@04$ with the model given in listing 4.4. In some cases the result might be one of the literal entries, meaning a value for the variable. But since it is a reference we get a `VarEntry` containing the identifier `Ref!val!0`. However, we want to turn it into a reference-entry containing all the fields, so we start going through the previously extracted heap in listing 4.2 and look for fields of the current identifier. We find one extracted chunk for the identifier `Ref!val!0`, storing the value of the field `next` given by another `VarEntry` containing `Ref!val!2`. We resolve

4. IMPLEMENTATION

it recursively. For the reference `Rev!val!2` we find two fields containing the literal integer-entries for 0 and 1 and we finally obtain the following `ExtractedModelEntry` for the value of `x`:

```
x <- Ref (Ref!val!0) {  
  next <- Ref (Ref!val!2) {  
    field1 <- 1  
    field2 <- 0  
  }  
}
```

While going through the fields of a nested reference, we keep a list of identifiers we already encountered. If we find a reference for a second time, instead of mapping its fields again, we return a `RecursiveRefEntry` to avoid infinite recursion if there are cyclic references. Finally, if we encounter a variable-entry of a sequence, we also resolve it to a `SequenceEntry` during the final mapping. If a counterexample involves a sequence, we will find the following two map-entries in the model:

```
Seq_length -> {  
  Seq<Int>!val!1 -> 5  
  else -> 0  
},  
Seq_index -> {  
  Seq<Int>!val!1 1 -> 4  
  Seq<Int>!val!1 2 -> 6  
  else -> 3  
}
```

Given the identifier from the variable-entry, in this case `Seq<Int>!val!1`, we can first get the length of the sequence from the `Seq_length` field. Then we can get the value at each index by evaluating the `Seq_index` function given the identifier and index as the arguments. We get a sequence of length 5 with the elements [3, 4, 6, 3, 3].

Recall that we always have a main heap, and a list of heaps for all the labels of the Viper program. After the mapping, our result is a main model, giving us the values of the method when the violation occurs, and a map from labels to older models. The map always contains a label "old", giving us the model for the values of variables when entering the method, and additional labels if the method contains explicitly defined labels.

4.2 Prusti Counterexamples

In this section we will explain how Silicon’s new representation of counterexamples can be used to generate counterexamples for Rust programs. To enable counterexample generation in our extension of Prusti, one has to set the environment variable `PRUSTI_COUNTEREXAMPLE` to `true`. In that case Prusti passes the additional flag “`-counterexample mapped`” to Silicon. If Silicon’s verification fails for a function, we should be able to get a counterexample for it. The produced counterexample is then copied over into Rust to a data structure equivalent to Silicon’s mapped counterexample. After we encounter a verification error and obtain a Silicon counterexample, we proceed with the following simplified steps: we identify the functions that fail verification, for each function we determine a set of Rust variables we are interested in, for each variable we look for its corresponding entry in the Silicon counterexample and translate it back to a Rust value, and finally we present the counterexample to the user. These steps will be explained in more detail in the following sections.

4.2.1 Identifying the failing function

Let us assume we are given an annotated Rust function that Prusti failed to verify and the counterexample option is enabled. Moreover, Silicon did produce a counterexample in the new representation. Before we can do any sort of processing, we need to know which functions of our Rust program failed. Therefore we need its so-called `DefId`, which uniquely identifies a function. To obtain the function’s id, we make use of the existing error-manager of Prusti. Its purpose is to map errors in Silicon back to the element in Rust that caused it, to report the failing statement. For every generated Viper-expression that might potentially cause a verification error, it registers an error location. They are registered during the encoding of the program, and we simply added the `DefId` of the currently encoded function to each of the registered errors. Whenever verification fails and Prusti identifies the reason for the failure, we are also given the identifier of the failing function.

4.2.2 Variables

After we find a failing function, we need to determine the variables that our counterexample should contain. We also want to know which of those variables are arguments of the function. Given the `DefId`, we can get the MIR-encoding of the function that fails verification. The variable names of the generated Viper program are the ones that we find in the MIR, however to give a counterexample we need to know which original Rust variables they correspond to. To find the variables of interest we use a data structure given by the compiler called `VarDebugInfo`, giving us a mapping from

source Rust variables to their corresponding MIR variables, for example `x -> _1`, `y -> _2`. They are always numbered starting from 1 and we are also given the number of arguments `arg_count`, telling us that the first `arg_count` variables are the function's arguments. Additionally, we know that the variable `_0` always stores the result of the function. Given the MIR identifier of a variable we can also get its type which means now, given a Silicon model we can try to get a value for this variable.

4.2.3 Backtranslation

In this section we look at how to translate an extracted model-entry back to a value of a Rust variable. For this part of the implementation, we need to be aware of the fact that a given model might not contain the entry for our variable at all or that its entry might not contain all the information we need, e.g. some fields of a reference might not be available. Our goal is to get as much information as possible from the given entries. The backtranslation is supported for variables of the following types: integers, booleans, chars, references, tuples, structs and enums. The result of the translation is a new type of entries, implemented as an enum with one variant for each of the supported types with the addition of unknown entries, which represent missing information or unsupported types. From now on we will always refer to Silicon's extracted model-entries as *Silicon entries*, and the ones we just introduced as *Rust entries*. The definition of the Rust entries is given in the following:

```
pub enum Entry {
    IntEntry { value: i64 },
    BoolEntry { value: bool },
    CharEntry { value: char },
    RefEntry { el: Box<Entry> },
    Struct {
        name: String,
        field_entries: Vec<(String, Entry)>,
    },
    Enum {
        super_name: String,
        name: Option<String>,
        field_entries: Vec<(String, Entry)>,
    },
    Tuple { fields: Vec<Entry> },
    Unit,
    UnknownEntry,
}
```

Assume we are given an identifier of a local MIR variable, its Rust type and an extracted Silicon model, and we want to get a value for the counterexample of the variable. Since the names of the MIR variables are the same as their translated Viper variables, we search the model for an entry of our current variable. If we find an entry we translate the found Silicon entry according to the given type-definition [3] of the Rust variable. Otherwise, we return an `UnknownEntry`. In the following sections, we discuss how the various Rust types are translated to Viper and how we can backtranslate a Silicon entry to a Rust entry given the type it encodes. More details on the Rust-Viper encoding performed by Prusti can be found in their main publication [4] or their developer guide [1].

Primitive Types

The encoding of integers, booleans and characters is straightforward. Any type of integer (unsigned, signed, long, short) in Rust is usually encoded to a reference with a field `val_int` in Viper. However, the generated encoding is sometimes simplified to just an integer. Depending on the Silicon entry we find, we either directly use the contained value, or if we get a reference we look for the field `val_int` and get its value. If we do not find the field we return an `UnknownEntry` and otherwise an `IntEntry` containing the found value. Two examples of given Silicon entries and the internal representation of the translated Rust entries are given in the following:

```
_1 <- 21
_2 <- Ref ($Ref!val!3) {
  val_int <- 43
}
```

Listing 4.5: Silicon

```
_1 <- IntEntry(21)
_2 <- IntEntry(43)
```

Listing 4.6: Prusti

The translation of booleans works analogously, with the sole difference that the field is named `val_bool` and we return a `BoolEntry`. For characters there is a difference because Viper does not have a `char` type. Rust characters are translated to integers in Viper, meaning we first get an integer value from the counterexample, as before. This value then has to be converted back to a `char` and we return a `CharEntry`. The following listings show an example of the translation of chars and booleans. Note that the entry for `_1` looks exactly the same as for an integer, but because the target variable in Rust is of type `char` it is translated differently:

```
_1 <- Ref {  
  val_int <- 63  
}  
_2 <- Ref {  
  val_bool <- true  
}
```

Listing 4.7: Silicon

```
_1 <- CharEntry('c'),  
_2 <- BoolEntry(true),
```

Listing 4.8: Prusti

References

Rust references are, not very surprisingly, encoded to Viper references. They have a field `val_ref` containing the value they point to. The provided type-definition also gives us the type of the value it points to. Given the Silicon entry generated for a Rust reference, we get its field `val_ref` and then recursively translate it according to the reference's sub-type. Then we return a `RefEntry` containing the result of the recursive translation. In the following listings, we show an example Silicon entry and the result of the translation for the type `Ref(i32)`:

```
_1 <- Ref {  
  val_ref <- Ref {  
    val_int <- 0  
  }  
}
```

Listing 4.9: Silicon

```
_1 <- RefEntry(IntEntry(0))
```

Listing 4.10: Prusti

Tuples

Tuples in Rust are also translated to references in Viper and their elements are stored in fields named `tuple_0`, `tuple_1`, etc. The type information gives us a list of the types of all the elements. Depending on the length of this list we try to find all the entries for those fields of the reference, and again translate the entries we find recursively. If we do not find a field, the difference to the previous types is that only the element that can not be found will be an `UnknownEntry`. Translating all the fields gives us a vector of Rust entries which we then return within a `TupleEntry`. In the example below we show a Silicon entry that is generated for a tuple of the type `Tuple(i32, bool)`. Note that the references pointed to by `tuple_0` and `tuple_1` have the same fields, which comes from the fact that they both point to the same reference.

```

_1 <- Ref {
  tuple_0 <- Ref {
    val_int <- 32
    val_bool <- false
  }
  tuple_1 <- Ref {
    val_int <- 32
    val_bool <- false
  }
}

```

Listing 4.11: Silicon

```

_1 <- TupleEntry([
  IntEntry(32),
  BoolEntry(false)
])

```

Listing 4.12: Prusti

Algebraic Data Types

Enums and structs in Rust belong to the same type called algebraic data types (ADT). Unions are also part of the same class but they are not supported by Prusti yet. We focus on structs first since they are just enums with only one variant in Rust. A struct is translated to a reference with a field for each field of the struct. The type's definition provides us with the name of the struct and all the names of the fields and their types, which lets us find the Silicon entries for the fields and evaluate them recursively again. Listing 4.13 shows an example of a struct definition in Rust.

Listing 4.13: Example of a struct definition

```

struct SomeStruct {
  value: i32,
  other_value: i32,
  valid: bool,
}

```

The Silicon entry for an instance of the given struct has the structure of the following example (left) and translating it gives us the shown Rust-entry (right).

4. IMPLEMENTATION

```
_0 <- Ref {
  f$valid <- Ref {
    val_bool <- true
    val_int <- 43
  }
  f$value <- Ref {
    val_int <- 42
  }
  f$other_value <- Ref {
    val_bool <- true
    val_int <- 43
  }
}
```

Listing 4.14: Silicon

```
_0 <- Struct {
  name: "SomeStruct"
  fields: [
    value: LitIntEntry(42),
    other_value: LitIntEntry(43),
    valid: BoolEntry(true)
  ]
}
```

Listing 4.15: Prusti

To extend this translation to enums, we will consider the following example definition:

```
enum Choose {
  One,
  Two{x: i32, y: bool},
  Three(char, bool)
}
```

To encode an enum, Prusti generates a reference with fields for each variant and an additional field storing the enum's discriminant. Each variant is encoded like a struct. The discriminant is an integer which has a certain value for each variant. An example of a Silicon entry that was generated for the above enum is shown in Listing 4.16

To translate the Silicon entry of an enum, we always try to get its discriminant first because it determines the variant of the enum. For our current example we get the value 2 for the discriminant. Given a discriminant and the internal representation of the enum's definition, we can get the name of the variant and the names of all its fields. In this case the discriminant tells us that the Silicon entry represents an instance of the variant `Three`. Therefore we have to get the entry of the field `enum_Three` and evaluate it as if it was a struct. For the given example we will encounter recursive references, which means we have to look for the complete reference entry with the same identifier to resolve some of the fields. Finally we get the following Rust entry:

Listing 4.16: Examples of a Silicon entry generated for enum Choose

```

_1 <- Ref ($Ref!val!0) {
  enum_Two <- Ref ($Ref!val!2) {
    f$0 <- RecRef($Ref!val!2)
    f$1 <- RecRef($Ref!val!2)
    val_int <- 99
    val_bool <- false
  }
  discriminant <- 2
  enum_Three <- Ref ($Ref!val!2) {
    f$0 <- RecRef($Ref!val!2)
    f$1 <- RecRef($Ref!val!2)
    val_int <- 99
    val_bool <- false
  }
}

```

Listing 4.17: translated Rust-entry

```

Enum {
  super_name: "Choose",
  name: "Three",
  field_entries: [
    0: CharEntry('c'),
    1: BoolEntry('false'),
  ]
}

```

4.2.4 Choosing Labels

At this point we have a set of Rust variables for which we want values in the counterexample, their corresponding MIR and Viper variable names and the capability to translate Silicon entries back to Rust entries. We now have to choose the models of the Silicon counterexample from which we want to get the values for the entries. To get the values of all variables at the end of the programs execution, the intuitive choice is to use the main model and for the initial values of the arguments the most obvious choice would be the model at the label "old".

However, after implementing this part of the project, we noticed that choosing the correct model is more subtle given the internal workings of Silicon. The problem is that in Viper, to access fields of references, one needs the permission to do so. In the programs generated by Prusti, most fields are accessed via predicates that are unfolded when those values are accessed

and folded again at a later point. At any label of the Viper program, if the predicate granting permission to the fields of a reference has not been previously unfolded or is already folded again, the counterexample does not know the fields values. To put this into context we consider the following Rust function that takes a struct (cf. Listing 4.13) as an argument:

```
pub fn foo(x: SomeStruct) {  
    assert!(x.value == x.other_value || x.valid)  
}
```

As we already know, the struct `x` will be translated to a reference `_1` with multiple fields. We will now look at a very reduced version of the generated Viper-method, containing only the statements that are relevant for whether or not fields of `_1` are accessible, which is shown in Listing 4.18.

Listing 4.18: predicate-definition labels and statements involving permissions of Viper program generated by Prusti for the Rust function in Section 4.2.4

```
1 predicate m_SomeStruct$_beg$_end_(self: Ref) {  
2     acc(self.f$value, write) && (acc(i32(self.f$value), write)  
3         && (acc(self.f$other_value, write)  
4             && (acc(i32(self.f$other_value), write)  
5                 && (acc(self.f$valid, write)  
6                     && acc(bool(self.f$valid), write))))))  
7 }  
8  
9 method m_foo() returns (_0: Ref)  
10 {  
11     var _1: Ref  
12     label start  
13     inhale acc(m_SomeStruct$_beg$_end_( _1), write)  
14  
15     label pre  
16     inhale true  
17     unfold acc(m_SomeStruct$_beg$_end_( _1), write)  
18     unfold acc(i32(_1.f$value), write)  
19  
20     label l0  
21     unfold acc(i32(_1.f$other_value), write)  
22  
23     label l1  
24     unfold acc(bool(_1.f$valid), write)  
25  
26     label l4  
27     ...  
28 }
```

`m_SomeStruct$.beg_$.end_` is the name of the predicate that gives access to the fields of the struct once it is unfolded. Since this unfolding (line 17) takes place after the label `pre` (line 15), these fields will only be in the counterexample for the models at the labels after it. The access to the fields of the actual integer and boolean values is unfolded later. Only at label 14 and after will all values of the struct be available. We can see this in the following entries of the generated Silicon counterexample at the various labels:

Listing 4.19: Silicon entries for `_1` of Listing 4.18 at the various labels

```
// at label old, start and pre:
_1 <- Ref ($Ref!val!0) {
}

// at label 10:
_1 <- Ref ($Ref!val!0) {
  f$valid(perm: 1/1) <- Ref ($Ref!val!3) {
    val_int(perm: 1/1) <- 0
  }
  f$other_value(perm: 1/1) <- Ref ($Ref!val!2) {
  }
  f$value(perm: 1/1) <- Ref ($Ref!val!3) {
    val_int(perm: 1/1) <- 0
  }
}

//at label 14:
_1 <- Ref ($Ref!val!0) {
  f$valid(perm: 1/1) <- Ref ($Ref!val!3) {
    val_bool(perm: 1/1) <- false
    val_int(perm: 1/1) <- 0
  }
  f$other_value(perm: 1/1) <- Ref ($Ref!val!2) {
    val_int(perm: 1/1) <- 1
  }
  f$value(perm: 1/1) <- Ref ($Ref!val!3) {
    val_bool(perm: 1/1) <- false
    val_int(perm: 1/1) <- 0
  }
}
}
```

In a similar fashion, if we want to get a value for the result of a function we can usually not use the main model because at this point the permissions for the fields of the result are already exhaled again. Unfortunately, there is no simple rule for where to evaluate variables such that we can be certain that we get a complete entry. In some cases, the unfolding and folding even

happens within the same label, such that the values can not be obtained at any label. To solve this problem we considered various solutions:

- when verification fails, insert additional fold/unfold statements at specific places and rerun verification to get more information;
- use the control flow graph of the program to find the labels where the values should be unfolded and evaluate them there; and
- change Silicon's implementation in a way that makes these values accessible even when they are not unfolded;

We dismissed the first idea because it would clutter the translation and double the running time. The second idea would still not always work, because in some cases the values are not available at any label as previously mentioned. Even though we might get more information than we currently do, we decided against this approach because it is more of a work-around than an actual solution to the problem. The third idea could potentially make all the values available at any label. Unfortunately the changes required in Silicon are quite involved and were not feasible in time for our project but they might be part of another project in the future.

For now, Prusti simply evaluates its arguments at the label 10 and all variables (including the arguments) at the last label that can be found in the list of models. Evaluating the arguments at the label 10 makes sense because at that point mutable arguments are usually not modified yet and we can often get some useful information already. We choose the last labelled model instead of the main model because the main model almost never contains the value for the result even though we might not get the true final value of a variable in some cases. The final result, after translating the given models as we have just described, gives us one or two Rust entries for each local variable and the result.

4.2.5 Presentation

The last section of this chapter describes how we present the counterexample to the user. Note that all the results of the backtranslation we have previously seen explicitly showed the internal representation of Rust entries and not how they are presented to the user. Their actual output is close to how they would be defined in Rust itself. For example the Rust-entry that was the result of the translation of a struct in Listing 4.15 is presented in the following way:

```
x <- SomeStruct {
  value: 0,
  other_value: 1,
  valid: false,
}
```

We will show more examples of different types during the evaluation in Chapter 5. As mentioned in Chapter 3, the entries are presented as compiler notes. For each variable we print a note containing the entries and its span. The compiler also automatically outputs the line of code where the variable is defined and underlines the variable. While the initial reason to use compiler notes was to make the output compatible with Prusti-Assistant, it has the additional advantage that we can easily identify variables in case of duplicate names. For example the function in Listing 4.20 generates the following two compiler notes for the two variables named `x`:

Listing 4.21: Prusti's presentation of counterexamples

```
note: counterexample for "x"
x <- 7
--> counterexample-thesis-resources/cetests/ref2.rs:10:9
|
10 |         x => x * 2
    |         ^
note: counterexample for "x"
initial: x <- ref(7)
final: x <- ref(7)
--> counterexample-thesis-resources/cetests/ref2.rs:5:8
|
5 | fn foo(x: &mut i32) -> i32{
  |         ^
```

Listing 4.20: Rust example with duplicate names

```
#[ensures(result != 14)]
fn foo(x: &i32) -> i32{
    let y = *x;
    match y {
        x => x * 2
    }
}

fn main(){}

```

4. IMPLEMENTATION




Figure 4.2: generated list of problems in VSCode when using Prusti-Assistant with our Version of Prusti and counterexamples enabled

From this output it is clear, which variable each note corresponds to. If we wanted to present the counterexample in a similar fashion to Nagini we would have to interpret the spans of variables in case of duplicate names. While the generated notes are currently not used by the VSCode extension Prusti-Assistant [5], it still recognizes and displays them in a more compact way as shown in Figure 4.2. Clicking a note makes the cursor jump directly to the variable it refers to. In the future this should eventually allow for a better IDE integration, where the variables would actually be underlined in the source code and hovering over them would display the generated counterexample.

Evaluation

The evaluation of this project consists of two parts: a qualitative evaluation of the generated counterexamples and an analysis of their influence on Prusti’s performance.

5.1 Qualitative Evaluation of Counterexamples

To evaluate the generated counterexamples, which are the main result of this thesis, we will look at a set of example Rust programs that fail verification in Prusti. Initially, our plan was to evaluate the quality of the counterexamples using Prusti’s existing tests, namely those that fail verification. Before starting the implementation, we annotated all of those tests with what we would consider optimal counterexamples. However, due to the limitations of our implementation caused by the often inaccessible values of counterex-

Example	Complete	Spurious	RE	LOC	VT	Source
fail-pc.rs	✓	✗	✓	10	0.49	Prusti tests [10]
replace.rs	✓	✗	✓	13	0.46	
tuple.rs	✗	✗	✓	8	0.71	
account.rs	✗	✗	✓	18	1.05	
enum.rs	✗	✗	✓	13	0.80	
account-fail.rs	✗	✗	✗	11	0.51	
loop.rs	✓	✓	✗	11	0.75	Prusti tests [10]
assert-bool.rs	✓	✗	✗	3	0.45	

Table 5.1: The list of examples used for the evaluation and a summary of the attributes of their generated counterexamples. We evaluated their completeness, i.e. the counterexamples contains no missing values, and whether the counterexample lets us reproduce the error (RE). It also contains the example’s lines of code (LOC) and verification time (VT) in seconds with counterexamples enabled.

amples (cf. Section 4.2.4), we decided to pick examples by hand that illustrate the strengths and weaknesses of our implementation. We evaluate the backtranslation for various types, discuss the validity of the generated counterexamples and then showcase limitations of the implementation. Table 5.1 summarizes the results.

Going back to our very first example of the introduction, shown again in Listing 5.1, we can see the generated output of the new Prusti version in Listing 5.2. This function is part of Prusti’s test suite and also shows our manual annotation of an “ideal counterexample” as we previously discussed. The generated counterexample is valid since it gives us a negative value for n and the program returns 0 for that value, which indeed violates the postcondition. Therefore this example does meet our initial expectations. The generated output is comparable to other compiler errors (e.g. invalid borrows) and thus the format is familiar to Rust programmers. Nonetheless, a more concise error message in an IDE (e.g. as we have seen in Figure 2.1 for Dafny) might be preferable. For the rest of the evaluation we will not be showing all of the generated output and reduce it to only the information about the variables.

Listing 5.1: fail-pc.rs (function from Prusti’s test suite with counterexample annotation)

```
#[ensures(result == n*(n+1)/2)]
fn sum(n: i32) -> i32 {
    if n <= 0 {
        0
    } else {
        sum(n-1) + n
    }
}
/* Counterexample:
   n <- -1
   result <- 0
*/
```

Listing 5.2: Prusti’s output for sum.rs

```
error: [Prusti: verification error] postcondition might not hold.
--> counterexample-thesis-resources/evaluation/sum.rs:5:11
|
5 | #[ensures(result == n*(n+1)/2)]
|   ~~~~~
note: the error originates here
--> counterexample-thesis-resources/evaluation/sum.rs:5:11
|
5 | #[ensures(result == n*(n+1)/2)]
|   ~~~~~
```



```

note: counterexample for "n"
initial: n <- -2
final: n <- -2
--> counterexample-thesis-resources/evaluation/sum.rs:6:8
|
6 | fn sum(n: i32) -> i32 {
|   ^
note: result <- 0
--> counterexample-thesis-resources/evaluation/sum.rs:6:19
|
6 | fn sum(n: i32) -> i32 {
|                   ^^^

Verification failed

```

Next we consider an example involving references, chars and booleans. Listing 5.3 contains a function that checks if the given reference points to the char '\$'. If it does, depending on whether the argument `acc` is true, it mutates the char to a white-space or panics. The generated counterexample, given in Listing 5.4, provides us with the correct arguments leading to a panic. It also illustrates how we decided to present chars. Apart from the char itself, we also displays its hexadecimal value. Additionally, we note that the counterexample gives us a value for the result even though the function would encounter a panic, an unrecoverable error in Rust, before returning any result. Our implementation currently does not detect whether the reason for the verification error is part of the function (e.g. assertion error, panic, etc.) or its post-condition and tries to get a value for the result in both cases, even though it should not be part of a counterexample in the former case.

Listing 5.3: `replace.rs`

```

fn replace(x: &mut char, acc: bool) {
    match x {
        '$' => {
            if acc {
                *x = ' ';
            } else {
                panic!("no_access");
            }
        },
        _ => {}
    }
}

```

5. EVALUATION

Listing 5.4: generated counterexample for replace.rs

```
initial:
  acc <- true
  x <- ref('%' (0x25))

final:
  acc <- true
  x <- ref('%' (0x25))
  result <- ()
```

The next function, shown in Listing 5.5, is given a tuple as an argument with the precondition that the first element of the tuple has to be larger than zero and the contained char has to be 'c'. It returns another tuple where the two entries of the argument are swapped and two is subtracted from the integer field. We try to prove that the second field of the result is always larger or equal than zero.

The resulting counterexample is shown in Listing 5.6. Looking at the initial value of the argument we can see that the first field of the tuple is available and correct since 1 is the only value that does not violate the precondition but leads to a violation of the postcondition. The question mark in the second field represents an `UnknownEntry` indicating that the corresponding Viper field has not been unfolded at label 10 where the initial values are

Listing 5.5: tuple.rs

```
#[requires(x.0 > 0 && x.1 == 'c')]
#[ensures(result.1 >= 0)]
fn foo(x: (i32, char)) -> (char, i32) {
  let y = x.0 - 2;
  let z = x.1;
  (z, y)
}
```

Listing 5.6: generated counterexample for tuple.rs

```
initial:
  x <- (1, ?)

final:
  x <- (1, 'c' (0x63))
  y <- -1
  z <- 'c' (63)
  result <- ('c' (0x63), -1)
```

evaluated. However, in the set of final values of the counterexample we can still find its value which is also correct. The result and the intermediate values also match an execution of the function and therefore this is also a valid counterexample.

For structs, we consider the function in Listing 5.7. The function `transfer` takes in two accounts and transfers an amount from account `x` to account `y`. The post-condition it tries to prove does not hold and finding a counterexample is trivial. However, this time we are handling mutable arguments and are interested in the values for `x` and `y` before and after they are mutated.

In the generated counterexample, given in Listing 5.8, we can see that most values are available. The initial amount and the initial balance of `y` is missing. The missing amount does not really pose a problem, since it is not mutated and must be equal to its initial value. We can see that the balance of `x` got reduced by the correct amount, but for `y` we could not be sure if we hadn't created the additional variable `temp` storing its value. Missing values are a major problem of our implementation, but in many cases we can still deduct what the reason for a failure was if we analyze the available values.

Listing 5.7: `account.rs`

```
pub struct Account {  
    balance: i32,  
}  
  
#[requires(amount > 0 && x.balance > amount && y.balance >= 0)]  
#[ensures(old(y.balance) > result.1.balance)]  
pub fn transfer(  
    mut x: Account,  
    mut y: Account,  
    amount: i32  
) -> (Account, Account) {  
    if x.balance >= amount {  
        let temp = y.balance;  
        x.balance -= amount;  
        y.balance += amount;  
    }  
    (x, y)  
}
```

Listing 5.8: generated counterexample for account.rs

```
initial:
  x <- Account { balance: 2701 }
  y <- Account { balance: ? }
  amount <- ?
  amount <- 2700

final:
  final: x <- Account { balance: 1 }
  y <- Account { balance: 5399 }
  temp <- 2699
  result <- (
    Account { balance: 1 }
    Account { balance: 5399 }
  )
```

Moving on to the last supported type, Listing 5.9 shows a function that evaluates enums representing three binary operations. The operation that is causing the error is `BinOp::Div` because it can lead to a division by zero. The generated counterexample, shown in Listing 5.10 does give us this value for the argument, but once again the fields are not available in the initial entry. The counterexample also contains all the variables for each arm of the match statement. This is unpleasant since many of those variables are never created during an execution which makes the counterexample wrong and less readable. Resolving this would require some analysis of the paths taken but for any analysis of this sort the counterexample for the arguments would have to be complete.

Listing 5.9: enum.rs

```
enum BinOp {
  Add(i32, i32),
  Sub(i32, i32),
  Div(i32, i32),
}

fn apply(op: BinOp) -> i32 {
  match op {
    BinOp::Add(a, b) => a + b,
    BinOp::Sub(c, d) => c - d,
    BinOp::Div(e, f) => e / f
  }
}
```

Listing 5.10: generated counterexample for Listing 5.9

```
initial:
  op <- BinOp::Div(?, ?)
final:
  op <- BinOp::Div(2, 0)
  a <- 0
  b <- 0
  c <- 0
  e <- 2
  f <- 0
  result <- 0
```

Listing 5.11: account-fail.rs

```
fn get_balance(acc: Account) -> i32 {
  acc.balance
}

#[ensures(result)]
fn has_money(acc: Account) -> bool {
  get_balance(acc) > 0
}
```

So far, even though we have mentioned the limitations of our approach several times, we have only shown examples where the inaccessible values of a counterexample were not actually relevant to trace the verification failure back to what causes it. However, in some cases we are not able to derive the reason for the verification failure from the counterexample and an example function for this behavior is given in Listing 5.11. The function `has_money` that fails verification operates on the same struct `Account` that was defined in Listing 5.7. The post-condition is violated for any account with a negative balance, but the important difference for this example is that the balance is accessed via an additional function. Because of that, the fields of the Viper

Listing 5.12: generated counterexample for Listing 5.11

```
initial:
  acc <- Account { balance: ? }
final:
  acc <- Account { balance: ? }
  result <- false
```

variable `_1`, the translated variable of `acc`, will never be unfolded in Viper and the value causing the failure is not accessible at any label of Silicon's counterexample. Therefore the generated counterexample, shown in Listing 5.12, does not contain any valuable information about the arguments.

Next we discuss an example of a spurious counterexample. Spurious behavior is expected in cases, where the reachable abstract states contain an error whereas the concrete state of this error is not reachable. Listing 5.13 shows an example where we expect this behavior. The function consists of a loop and a post-condition that will not be violated when executing the function. However, the loop invariant is too weak to prove that the assertion holds and therefore verification fails. After symbolic execution of the loop the verifier knows that the equation `y > 0 && x >= 0 && !(x > 0)` holds, which is an over-approximation of the concretely reachable states. The generated counterexample gives us a final value of 16 for `y` which is what we expected. However, if we replace the the post-condition with an assertion `assert!(y < 16)` at the end of the function, Prusti generates 15 for the value of `y` which does not violate the assertion. Even though both of the results are not valid counterexamples, the second result is not expected and is probably caused by some implementation error.

Listing 5.13: `loop.rs`

```
#[ensures(result < 16)]
fn spurious() -> i32 {
    let mut x = 10;
    let mut y = 1;
    while(x > 0) {
        body_invariant!(x >= 0 && y > 0);
        x = x - 1;
        y = y + 1;
    }
    y
}
```

At this point we encountered several examples involving assertions that resulted in counterexamples that would not lead to a violation. The most blatantly incorrect example is shown in Listing 5.14 where the generated counterexample is `texttt{x} - true`. However, after investigating the source of the error we determined that Silicon is responsible for this value. The counterexample Silicon generates is also wrong for the translated Viper program and the error was not introduced by our extraction in Silicon.

Listing 5.14: assert-bool.rs

```
fn foo(x: bool) {
    assert!(x);
}
```

5.2 Timing Analysis

For the evaluation of the influence on Prusti’s performance when generating counterexamples, we adjusted Prusti’s existing benchmarking script to run a set of 121 Rust files that fail verification with counterexamples enabled and disabled. The tests consist of the examples used during the evaluation and Prusti’s existing tests that fail verification.

The results of this analysis showed that verification with counterexamples enabled takes on average around 0.1 seconds longer than when they are disabled. The distribution of the increase of the execution time in seconds is shown in Figure 5.1. In most cases the difference is less than 0.2 seconds but there are a few outliers where the difference is quite substantial. Analyzing the extreme cases where the increase is more than 2 seconds showed that both of them involved loops, however other examples involving loops did not confirm a trend of high verification times for loops in general.

Additionally, we were interested in how this increase is related to the total verification time. The results show that enabling counterexamples increases the verification time on average around 6.5%. Figure 5.2 shows that the influence can be quite substantial in some cases, but is smaller than 20% in a majority (>90%) of the tests.




Figure 5.1: Histogram for increase in verification-time when using counterexamples

5. EVALUATION




Figure 5.2: Histogram for relative increase in verification-time

The influence on the performance should be barely noticeable in most cases and should not have an influence on the verification experience. Since counterexamples will mostly be used as a debugging tool and not be constantly enabled, this seems acceptable.

Chapter 6

Conclusion

Appendix A

Appendix

Bibliography

- [1] Prusti Developer-Guide on Type Encoding. <https://viperproject.github.io/prusti-dev/dev-guide>. Accessed: 02.05.2021.
- [2] The Rust programming language. <https://www.rust-lang.org>. Accessed: 28.04.2021.
- [3] Rust Type Definitions. https://doc.rust-lang.org/stable/nightly-rustc/rustc_middle/ty/sty/enum.TyKind.html. Accessed: 02.05.2021.
- [4] V. Astrauskas, P. Müller, F. Poli, and A. J. Summers. Leveraging rust types for modular specification and verification. *Proc. ACM Program. Lang.*, 3(OOPSLA):147:1–147:30, 2019.
- [5] J. Dunskus. Developing IDE Support for a Rust Verifier. Master’s thesis, ETH Zürich, 2020.
- [6] M. Eilers and P. Müller. Nagini: A static verifier for python. In *CAV (1)*, volume 10981 of *Lecture Notes in Computer Science*, pages 596–603. Springer, 2018.
- [7] K. Rustan M. Leino. Dafny: An automatic program verifier for functional correctness. In *LPAR (Dakar)*, volume 6355 of *Lecture Notes in Computer Science*, pages 348–370. Springer, 2010.
- [8] P. Müller, M. Schwerhoff, and A. J. Summers. Viper: A verification infrastructure for permission-based reasoning. In B. Jobstmann and K. R. M. Leino, editors, *Verification, Model Checking, and Abstract Interpretation (VMCAI)*, volume 9583 of *LNCS*, pages 41–62. Springer-Verlag, 2016.

BIBLIOGRAPHY

- [9] C. Stoll. Smt models for verification debugging. Master's thesis, ETH Zürich, 2019.
- [10] Prusti team. Prusti source code. <https://github.com/viperproject/prusti-dev/>. Accessed: 28.04.2021.



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Declaration of originality

The signed declaration of originality is a component of every semester paper, Bachelor's thesis, Master's thesis and any other degree paper undertaken during the course of studies, including the respective electronic versions.

Lecturers may also require a declaration of originality for other written papers compiled for their courses.

I hereby confirm that I am the sole author of the written work here enclosed and that I have compiled it in my own words. Parts excepted are corrections of form and content by the supervisor.

Title of work (in block letters):

Authored by (in block letters):

For papers written by groups the names of all authors are required.

Name(s):

First name(s):

With my signature I confirm that

- I have committed none of the forms of plagiarism described in the '[Citation etiquette](#)' information sheet.
- I have documented all methods, data and processes truthfully.
- I have not manipulated any data.
- I have mentioned all persons who were significant facilitators of the work.

I am aware that the work may be screened electronically for plagiarism.

Place, date

Signature(s)

For papers written by groups the names of all authors are required. Their signatures collectively guarantee the entire content of the written paper.